PLAYBOOK . CHATGPT

A first *safe week* for people who use ChatGPT at work

ChatGPT can be a personal account or one your employer provides. The first question is whose account it is.

If you are not sure where to start

People often do not call it fear. Said out loud, it sounds like one of these:

- "**What if I break something?**" The first week asks you to send nothing, delete nothing, and put nothing from your job in.
- "**I do not know what to do.**" Each day has one small job, and the first one is reading a document.
- "**I do not know where to start.**" Start with the rules at your company. That is Monday.
- "**I do not know what it could do for me.**" The section on where people find real time shows the places worth trying.
- "**What if I cannot do what I am supposed to do with it?**" You do not need to be technical. You need to check the work, and the kit gives you a way to do that.
- "**Am I even allowed?**" That is the policy check, and it comes first.

Each of these is a fair thing to wonder, and each one gets a next step below.

This guide works with any tool. Copilot is often the easy one to switch on, since many large companies already run on Microsoft, but it is not the only one. The playbooks for ChatGPT, Gemini, and Claude cover the others.

> **Before anything else: check your employer's policy**
>
> Your company decides what you may put into an AI tool, and which tools you may use. Find the written policy (often called an acceptable use, AI use, or information security policy). If you cannot find one, ask your manager or your IT or security team, and keep the answer in writing. Everything in this guide sits under that policy. If the two ever disagree, the policy wins.

ChatGPT call-outs

- **Whose account is it?** A personal account is not your employer's tool. If your company provides a work workspace, use that for work. If it does not, ask before using a personal account for anything from work.
- **Custom Instructions** are the right home for the context note. OpenAI says they apply to all your chats, immediately, and are available on all plans.
- **Memory is a separate layer.** OpenAI describes saved memories, which you can ask ChatGPT to forget or delete in Memory settings, and a "reference chat history" option that can change over time. Deleting a saved memory does not remove mentions of it from past conversations.
- **Temporary chats** do not create or update memories, per OpenAI, even with personalization on. Use one for Step 1 practice.
- **Workspace plans.** OpenAI says that by default it does not use content from ChatGPT Business, Enterprise, or Edu workspaces to train its models. Which plan you are on is something to confirm with your IT team.

Sources, read on 2026-10-17. OpenAI Help Center, "Memory in ChatGPT": help.openai.com/en/articles/8590148-memory-in-chatgpt. OpenAI Help Center, "ChatGPT Custom Instructions": help.openai.com/en/articles/8096356-chatgpt-custom-instructions. Product settings change, so check your own screen.

Not affiliated with or endorsed by OpenAI. ChatGPT is OpenAI's name.

The risk ladder

Not every use of AI carries the same risk. Start at the bottom of the ladder and move up only when you have a reason and the policy allows it.

Step 1 No company information goes in	Explain a concept you are new to. Brainstorm names, agendas, or questions for a meeting. Practice a hard conversation using made-up details. Ask how a spreadsheet formula works, using made-up rows. Draft a blank template you will fill in yourself. Turn your own rough notes about a public topic into an outline.
Step 2 Work content goes in, with care	Summarize your own meeting notes. Turn your bullets into a draft email. Tidy a list. Ask for the formula that fits your column names. Plan a dashboard or report. Only after the policy says yes, only the minimum content, names and numbers removed or replaced, and everything checked by you before it goes anywhere.

Step 3 Do not use AI for this	Personal data about employees, customers, patients, or candidates. Anything confidential, privileged, or not yet public. Passwords and keys. Anything your policy forbids. Decisions about a person's job, pay, or performance. Anything you have no way to check. Sending to a customer or a regulator without you reading every line.

Set your context once

A short plain-text note about your work saves you from retyping the same background in every chat, and keeps answers closer to what you need. Keep it as a file you own. Paste it at the start of a new chat, or into the custom instructions of the tool your company provides. Tools change; the file stays yours.

Never put anything in it that your policy would not let you type into the tool. Roles and reader types are fine. Names, numbers, and client details are not.

Where it goes: In ChatGPT, use Custom Instructions for the note, or paste it at the start of a new chat.

```
My role: [your job title, in a few words]
Who reads my work: [for example: my manager, the finance team, customers]
How I write: [short sentences, plain words, no jargon]
Words I use: [terms people in my team use]
What stays out: [anything I never put in a chat]
How to answer: ask me one question if something is unclear, say what you are
unsure of, and show your steps when you do math.
```

The verification habit

AI can be fluent and wrong in the same sentence. A spreadsheet with a broken formula opens exactly like a correct one, so you build the check yourself. Before anything leaves your hands:

1. **Numbers by hand.** Pick three rows and work them out yourself. If your math and the formula agree, keep it. If not, you caught it.
2. **Names, dates, and quotes against the original.** Open the source and compare. Do not ask the AI to confirm itself.
3. **Ask what it is unsure of.** "What in this answer might be wrong, and how would I check it?" is a fair question.
4. **Read it as the person who receives it.** Would they act on it? Would they trust it?
5. **When a draft fails twice, stop editing it.** Open a new chat with a bare brief: who reads it, what they should do, three things it must not do.

Your first safe week

Each day is a small step. The week asks you to send nothing, delete nothing, and put nothing from your job in. 15 to 20 minutes a day is plenty.

Monday	**Read the rules.** Find your employer's policy. Note which tool your company provides and sign in with your work account, not a personal one. Write down anything you are unsure about and ask.
Tuesday	**Step 1 only.** Use the tool for something with no company information in it: explain a concept, plan a meeting agenda with made-up details, ask how a formula works. Get a feel for how it answers.
Wednesday	**Set your context.** Fill in the context note above and paste it into a new chat. Ask for help with a task and notice how the answer changes.
Thursday	**Practice on a made-up or public example.** Pick a task that looks like your work but uses a pretend scenario or public information: a made-up email to a made-up colleague, or a public article to summarize. Run the verification habit on the result before you use any of it.
Friday	**Look back.** Three lines in your notes: what saved time, what needed fixing, what you will not use it for. Then fill in the first rows of the AI audit worksheet.

From week two, if your policy allows it: one small Step 2 task from your own work. Take out names and numbers first, put in only the minimum, and run the verification habit before you use any of it.

Beyond email: where people find the real time

Email is where most people start and stop. These are the places that tend to repay the effort, each with the same rule: you check the result.

- **Spreadsheet formulas.** Paste the column names, never the data. Add a few made-up rows with the same shape. Ask for the formula and a plain explanation of each piece. Check three rows by hand.
- **Report and dashboard planning.** Write down the decisions the report must support, then ask for the smallest set of views that serve them. Every tile should name the decision it serves.
- **Meeting preparation.** Questions to ask, objections to expect, and the one risk to raise, from your own agenda.

- **Plain-language rewrites.** Turn a dense policy or procedure you own into something a new hire can follow.
- **Practice conversations.** Rehearse a hard conversation against a made-up person before the real one.
- **Checklists and templates.** Draft the blank version of anything you do repeatedly, then fill it in yourself.
- **Reading long documents.** Ask for the decision a document is asking you to make, then go and read the parts that matter.
- **Finding what you have never defined.** Before asking AI to analyze your data, ask it what terms in your question need a definition.

Where AI should not be used

- **Anything your policy does not allow.** This beats every other rule in this guide.
- **Personal data about people.** Employees, customers, patients, candidates, students.
- **Anything confidential or not yet public.** Financial results, legal matters, contracts, plans, source code, client material.
- **Decisions about a person.** Hiring, firing, pay, ratings, discipline. A tool can help you prepare. The decision and the reasons are yours.
- **Passwords, keys, and access details.**
- **Anything you cannot check.** If you cannot tell whether the answer is right, it is not ready to use.
- **Work sent to someone else unread.** You are the author of whatever leaves under your name.

> **You do not have to know everything first.** Worrying you will break something, not knowing what to do or where to start, not knowing what AI could do for your work, or doubting that you can do what you are supposed to do with it: these are normal places to begin, and they are good instincts, not weaknesses. This guide gives each one a next step: a policy check first, a ladder you climb slowly, and a habit that catches errors before anyone else does. Starting small and checking as you go is a perfectly good way to learn this.
>
> General information for working people. Not legal, compliance, or HR advice. Your employer's policy and your own judgment decide what is right for you.